



DHAANISH AHMED COLLEGE OF ENGINEERING

Dhaanish Nagar, Padappai, Chennai – 601301

Approved By AICTE, New Delhi,

Affiliated to Anna University, Chennai.

www.dhaanish.in

Department of Artificial Intelligence & Data Science

Lab Manual

CCS334 – Big Data Analytics Laboratory

Year/Sem : III/V

DHAANISH AHMED COLLEGE OF ENGINEERING

Vision

To establish a world-class institution that is recognized as a “Centre of Excellence” offering education and research in engineering, technology and management with a blend of social and moral values to serve the community with a futuristic perspective.

Mission

To produce eminent engineers and managers with academic excellence in their chosen fields, which would be able to take up the challenges in the modern era and fulfill the expectations of the organization they join, with moral values and social ethics.

Department of Artificial Intelligence and Data Science

Vision

To impart quality Education, Industry Collaboration, promote Research and produce Graduate Industry-ready Engineers in the field of Artificial Intelligence and Data Science to serve the society.

Mission

- To provide a conducive learning environment for quality education in the field of Artificial Intelligence and Data Science.
- To promote industry-institute interaction and collaborative research activities.
- To empower the students with ethical values and social responsibilities in their profession.

PROGRAMME EDUCATIONAL OBJECTIVES (PEOs)

- Show proficiency in the knowledge of basic sciences, mathematics, Artificial Intelligence, data science and statistics to build systems that require management and analysis of large volume of data.
- Demonstrate technical skills to pursue pioneering research in the field of AI and Data Science and create disruptive and sustainable solutions for the welfare of ecosystems.
- Exhibit effective communication skills, team work and lead their profession with ethics.

Program Specific Outcome (PSO)

PSO1: Evolve AI based efficient domain specific processes for Effective decision making in several domains such as business and governance domains.

PSO2: Create, select and apply the theoretical knowledge of AI and Analytics along with practical industrial tools and techniques to manage and solve societal problems.

INDEX

S.NO	DATE	EXPERIMENT NAME	PAGE NO	SIGN
1		Install Apache Hadoop		
2		Develop a MapReduce program to implement Matrix Multiplication.		
3		Develop a MapReduce program to find the grades of student's.		
4		Develop a MapReduce program to calculate the frequency of a given word in agiven file.		
5		Develop a MapReduce to find the maximum electrical consumption in each year given electrical consumption for each month in each year.		
6		Develop a MapReduce to analyze weather data set and print whether the day is shinny or cool day.		
7		XYZ.com is an online music website where users listen to various tracks, the data gets collected .		
8		Develop a MapReduce program to analyze Uber data set to find the days on which each basement has more trips using the following dataset.		
9		Hive Operations		
10		Write queries to sort and aggregate the data in a table using HiveQL.		

Ex.No:01

Date:

Install Apache Hadoop

Aim:

To Install Apache Hadoop.

Hadoop software can be installed in three modes of

Hadoop is a Java-based programming framework that supports the processing and storage of extremely large datasets on a cluster of inexpensive machines. It was the first major open source project in the big data playing field and is sponsored by the Apache Software Foundation.

Hadoop-2.7.3 is comprised of four main layers:

□□ **Hadoop Common** is the collection of utilities and libraries that support other Hadoop modules.

□□ **HDFS**, which stands for Hadoop Distributed File System, is responsible for persisting data to disk.

□□ **YARN**, short for Yet Another Resource Negotiator, is the "operating system" for HDFS.

□□ **MapReduce** is the original processing model for Hadoop clusters. It distributes work within the cluster or map, then organizes and reduces the results from the nodes into a response to a query. Many other processing models are available for the 2.x version of Hadoop.

Hadoop clusters are relatively complex to set up, so the project includes a stand-alone mode which is suitable for learning about Hadoop, performing simple operations, and debugging

Algorithm:

1. Install Apache Hadoop 2.2.0 in Microsoft Windows OS

If Apache Hadoop 2.2.0 is not already installed then follow the post Build, Install, Configure and Run Apache Hadoop 2.2.0 in Microsoft Windows OS.

2. Start HDFS (Namenode and Datanode) and YARN (Resource Manager and Node Manager) .

Program:

Run following commands. *Command Prompt*

```
C:\Users\abhijitg>cd c:\hadoop
```

```
c:\hadoop>sbin\start-dfs
```

```
c:\hadoop>sbin\start-yarn
```

starting yarn daemons

Namenode, **Datanode**, **Resource Manager** and **Node Manager** will be started in few minutes and ready to execute Hadoop **MapReduce** job in the Single Node (pseudo-distributed mode) cluster.

```

Apache Hadoop Distribution - hadoop namenode
. infoPort=50075, ipcPort=50020, storageInfo=lv=-47;cid=CID-1af...d9f-efee-4d4e-9f03-...2c23e5eb;nsid=28920020;c=0> storage DS-2081780230-...-50010-1383498502153
13/11/03 22:38:22 INFO net.NetworkTopology: Adding a new node: /default-rack/127.0.0.1:50010
13/11/03 22:38:22 INFO blockmanagement.BlockManager: BLOCK* processReport: Received first block report from 127.0.0.1:50010 after starting up or becoming active. Its block contents are no longer considered stale
13/11/03 22:38:22 INFO BlockStateChange: BLOCK* processReport: from DatanodeRegistration(127.0.0.1, storageID=DS-2081780230-...-50010-1383498502153, infoPort=50075, ipcPort=50020, storageInfo=lv=-47;cid=CID-1af...d9f-efee-4d4e-9f03-...a0...b;nsid=28920020;c=0), blocks: 0, processing time: 15 msec

```

```

Apache Hadoop Distribution - hadoop datanode
207-50010-1383498502153> service to localhost/127.0.0.1:9000
13/11/03 22:38:22 INFO datanode.DataNode: BlockReport of 0 blocks took 0 msec to generate and 94 msec for RPC and NN processing
13/11/03 22:38:22 INFO datanode.DataNode: sent block report, processed command:org.apache.hadoop.hdfs.server.protocol.FinalizeCommand@78a6195e
13/11/03 22:38:22 INFO util.GSet: Computing capacity for map BlockMap
13/11/03 22:38:22 INFO util.GSet: UM type = 64-bit
13/11/03 22:38:22 INFO util.GSet: 0.5% max memory = 888.9 MB
13/11/03 22:38:22 INFO util.GSet: capacity = 2^19 = 524288 entries
13/11/03 22:38:22 INFO datanode.BlockPool$SliceScanner: Periodic Block Verification

```

```

Apache Hadoop Distribution - yarn resourcemanager
oop.yarn.server.api.ResourceManagerAdministrationProtocolPB to the server
13/11/03 22:48:14 INFO ipc.Server: IPC Server Responder: starting
13/11/03 22:48:14 INFO ipc.Server: IPC Server listener on 8033: starting
13/11/03 22:48:14 INFO util.RackResolver: Resolved ABHIJITG...com to /default-rack
13/11/03 22:48:14 INFO resourcemanager.ResourceTrackerService: NodeManager from node ABHIJITG...com(cmPort: 60092 httpPort: 8042) registered with capability: <memory:8192, vCores:8>, assigned nodeId ABHIJITG...com:60092
13/11/03 22:48:14 INFO rmnode.RMNodeImpl: ABHIJITG...com:60092 Node Transitioned from NEW to RUNNING
13/11/03 22:48:14 INFO capacity.CapacityScheduler: Added node ABHIJITG...com:60092 clusterResource: <memory:8192, vCores:8>

```

```

Apache Hadoop Distribution - yarn nodemanager
13/11/03 22:48:13 INFO mortbay.log: Started SelectChannelConnector@0.0.0.0:8042
13/11/03 22:48:13 INFO webapp.WebApps: Web app /node started at 8042
13/11/03 22:48:14 INFO webapp.WebApps: Registered webapp guice modules
13/11/03 22:48:14 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0.0:8031
13/11/03 22:48:14 INFO security.NMContainerTokenSecretManager: Rolling master-key for container-tokens, got key with id 441918079
13/11/03 22:48:14 INFO security.NMTokenSecretManagerInNM: Rolling master-key for nm-tokens, got key with id :1221761938
13/11/03 22:48:14 INFO nodemanager.NodeStatusUpdaterImpl: Registered with ResourceManager as ABHIJITG...com:60092 with total resource of <memory:8192, vCores:8>
13/11/03 22:48:14 INFO nodemanager.NodeStatusUpdaterImpl: Notifying ContainerManager to unblock new container-requests

```

Run wordcount MapReduce job

Now we'll run **wordcount** MapReduce job available in

%HADOOP_HOME%\share\hadoop\mapreduce\hadoop-mapreduce-examples-2.2.0.jar

Create a text file with some content. We'll pass this file as input to the **wordcount** MapReduce job for counting words. *C:\file1.txt*

Install Hadoop

Run Hadoop Wordcount Mapreduce Example

Create a directory (say 'input') in HDFS to keep all the text files (say 'file1.txt') to be used for counting words.

C:\Users\abhijitg>cd c:\hadoop

C:\hadoop>bin\hdfs dfs -mkdir input

Copy the text file(say 'file1.txt') from local disk to the newly created 'input' directory in HDFS.

```
C:\hadoop>bin\hdfs dfs -copyFromLocal c:/file1.txt input
```

Check content of the copied file.

```
C:\hadoop>hdfs dfs -ls input
```

Found 1 items

```
-rw-r--r-- 1 ABHIJITG supergroup 55 2014-02-03 13:19 input/file1.txt
```

```
C:\hadoop>bin\hdfs dfs -cat input/file1.txt
```

Install Hadoop

Run Hadoop Wordcount Mapreduce Example

Run the wordcount MapReduce job provided in

```
%HADOOP_HOME%\share\hadoop\mapreduce\hadoop-mapreduce-examples-2.2.0.jar
```

```
C:\hadoop>bin\yarn jar share/hadoop/mapreduce/hadoop-mapreduce-examples-2.2.0.jar  
wordcount input output
```

```
14/02/03 13:22:02 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0.0:8032
```

```
14/02/03 13:22:03 INFO input.FileInputFormat: Total input paths to process : 1
```

```
14/02/03 13:22:03 INFO mapreduce.JobSubmitter: number of splits:1
```

```
:
```

```
:
```

```
14/02/03 13:22:04 INFO mapreduce.JobSubmitter: Submitting tokens for job:
```

```
job_1391412385921_0002
```

```
14/02/03 13:22:04 INFO impl.YarnClientImpl: Submitted application
```

```
application_1391412385921_0002 to ResourceManager at /0.0.0.0:8032
```

```
14/02/03 13:22:04 INFO mapreduce.Job: The url to track the job:
```

```
http://ABHIJITG:8088/proxy/application_1391412385921_0002/
```

```
14/02/03 13:22:04 INFO mapreduce.Job: Running job: job_1391412385921_0002
```

```
14/02/03 13:22:14 INFO mapreduce.Job: Job job_1391412385921_0002 running in uber mode :  
false
```

```
14/02/03 13:22:14 INFO mapreduce.Job: map 0% reduce 0%
```

```
14/02/03 13:22:22 INFO mapreduce.Job: map 100% reduce 0%
```

```
14/02/03 13:22:30 INFO mapreduce.Job: map 100% reduce 100%
```

```
14/02/03 13:22:30 INFO mapreduce.Job: Job job_1391412385921_0002 completed successfully
```

```
14/02/03 13:22:31 INFO mapreduce.Job: Counters: 43
```

File System Counters

```
FILE: Number of bytes read=89
```

```
FILE: Number of bytes written=160142
```

```
FILE: Number of read operations=0
```

```
FILE: Number of large read operations=0
```

```
FILE: Number of write operations=0
```

```
HDFS: Number of bytes read=171
```

```
HDFS: Number of bytes written=59
```

```
HDFS: Number of read operations=6
```

```
HDFS: Number of large read operations=0
```

```
HDFS: Number of write operations=2
```

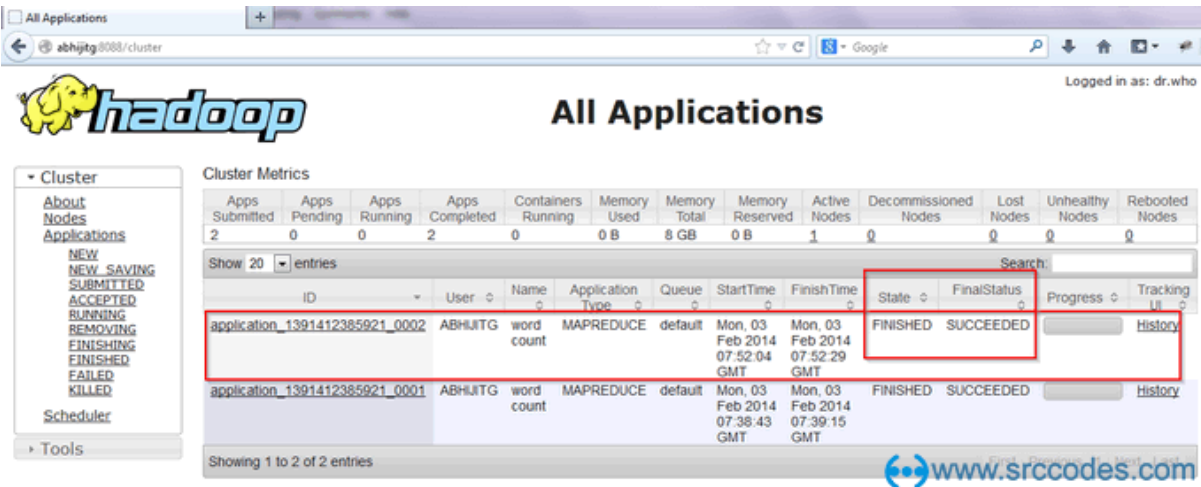
Job Counters

```
Launched map tasks=1
```

```
Launched reduce tasks=1
```

Data-local map tasks=1
Total time spent by all maps in occupied slots (ms)=5657
Total time spent by all reduces in occupied slots (ms)=6128
Map-Reduce Framework
Map input records=2
Map output records=7
Map output bytes=82
Map output materialized bytes=89
Input split bytes=116
Combine input records=7
Combine output records=6
Reduce input groups=6
Reduce shuffle bytes=89
Reduce input records=6
Reduce output records=6
Spilled Records=12
Shuffled Maps =1
Failed Shuffles=0
Merged Map outputs=1
GC time elapsed (ms)=145
CPU time spent (ms)=1418
Physical memory (bytes) snapshot=368246784
Virtual memory (bytes) snapshot=513716224
Total committed heap usage (bytes)=307757056
Shuffle Errors
BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0
File Input Format Counters
Bytes Read=55
File Output Format Counters
Bytes Written=59

Output:



The screenshot shows the Hadoop All Applications web interface. The browser address bar shows 'abhiitg@0083/cluster'. The page title is 'All Applications'. The user is logged in as 'dr.who'. The interface includes a sidebar with navigation links like 'Cluster', 'About', 'Nodes', 'Applications', 'NEW', 'NEW_SAVING', 'SUBMITTED', 'ACCEPTED', 'RUNNING', 'REMOVING', 'FINISHING', 'FINISHED', 'FAILED', 'KILLED', 'Scheduler', and 'Tools'. The main content area displays 'Cluster Metrics' and a table of applications. Two applications are listed, both in a 'FINISHED' state with a 'SUCCEEDED' final status. A red box highlights the 'State' and 'FinalStatus' columns for both entries.

ID	User	Name	Application Type	Queue	StartTime	FinishTime	State	FinalStatus	Progress	Tracking UI
application_1391412385921_0002	ABHIITG	word count	MAPREDUCE	default	Mon, 03 Feb 2014 07:52:04 GMT	Mon, 03 Feb 2014 07:52:29 GMT	FINISHED	SUCCEEDED		History
application_1391412385921_0001	ABHIITG	word count	MAPREDUCE	default	Mon, 03 Feb 2014 07:38:43 GMT	Mon, 03 Feb 2014 07:39:15 GMT	FINISHED	SUCCEEDED		History

Showing 1 to 2 of 2 entries

www.srccodes.com

Result:

We've installed Hadoop in stand-alone mode and verified it by running an example program it provided.


EXP.NO:02

DATE: Implement of Matrix Multiplication with Hadoop Map Reduce

AIM:

To Develop a MapReduce program to implement Matrix Multiplication.

In **mathematics**, **matrix multiplication** or the **matrix product** is a binary operation that produces a matrix from two matrices. The definition is motivated by linear equations and linear transformations on vectors, which have numerous applications in applied mathematics, physics, and engineering. In more detail, if **A** is an $n \times m$ matrix and **B** is an $m \times p$ matrix, their matrix product **AB** is an $n \times p$ matrix, in which the m entries across a row of **A** are multiplied with the m entries down a column of **B** and summed to produce an entry of **AB**. When two linear transformations are represented by matrices, then the matrix product represents the composition of the two transformations.



Algorithm for Map Function:

- a. for each element m_{ij} of M do produce (key,value) pairs as $((i,k), (M,j,m_{ij}))$, for $k=1,2,3,\dots$ upto the number of columns of N
- b. for each element n_{jk} of N do produce (key,value) pairs as $((i,k),(N,j,n_{jk}))$, for $i = 1,2,3,\dots$ Upto the number of rows of M .
- c. return Set of (key,value) pairs that each key (i,k) , has list with values (M,j,m_{ij}) and (N, j,n_{jk}) for all possible values of j .

Algorithm for Reduce Function:

- d. for each key (i,k) do
- e. sort values begin with M by j in list M sort values begin with N by j in list N multiply m_{ij} and n_{jk} for j th value of each list
- f. sum up $m_{ij} \times n_{jk}$ return $(i,k), \sum_{j=1} m_{ij} \times n_{jk}$

Step 1. Download the hadoop jar files with these links.

Download Hadoop Common Jar files: <https://goo.gl/G4MyHp>

\$ wget <https://goo.gl/G4MyHp> -O hadoop-common-2.2.0.jar

Download Hadoop Mapreduce Jar File: <https://goo.gl/KT8yfB>

\$ wget <https://goo.gl/KT8yfB> -O hadoop-mapreduce-client-core-2.7.1.jar

Step 2. Creating Mapper file for Matrix Multiplication.

```
import java.io.DataInput;
import java.io.DataOutput;
import java.io.IOException;
import java.util.ArrayList;
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.DoubleWritable;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.io.Writable;
import org.apache.hadoop.io.WritableComparable;
import org.apache.hadoop.mapreduce.Job;

import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.lib.input.*;
import org.apache.hadoop.mapreduce.lib.output.*;
import org.apache.hadoop.util.ReflectionUtils;
class Element implements Writable {
int tag;
int index;
double value;
Element() {
tag = 0;
```

```
index = 0;
value = 0.0;
}
Element(int tag, int index, double value) {
this.tag = tag;
this.index = index;
this.value = value;
}
@Override
public void readFields(DataInput input) throws IOException {
tag = input.readInt();
index = input.readInt();
value = input.readDouble();
}
@Override
public void write(DataOutput output) throws IOException {
output.writeInt(tag);
output.writeInt(index);
output.writeDouble(value);
}
}
class Pair implements WritableComparable<Pair> {
int i;
int j;
Pair() {
i = 0;

j = 0;
}
Pair(int i, int j) {
this.i = i;
this.j = j;
}
@Override
public void readFields(DataInput input) throws IOException {
i = input.readInt();
j = input.readInt();
}
@Override
public void write(DataOutput output) throws IOException {
output.writeInt(i);
output.writeInt(j);
}
@Override
public int compareTo(Pair compare) {
if (i > compare.i) {
```

```

return 1;
} else if ( i < compare.i) {
return -1;
} else {
if(j > compare.j) {
return 1;
} else if (j < compare.j) {
return -1;
}
}
return 0;
}
public String toString() {
return i + " " + j + " ";
}
}
public class Multiply
{
public static class MatriceMapperM extends Mapper<Object,Text,IntWritable,Element>

```

{ 24 Department of CSE

@Override

```

public void map(Object key, Text value, Context context)
throws IOException, InterruptedException {
String readLine = value.toString();
String[] stringTokens = readLine.split(",");
int index = Integer.parseInt(stringTokens[0]);
double elementValue = Double.parseDouble(stringTokens[2]);
Element e = new Element(0, index, elementValue);
IntWritable keyValue = new IntWritable(Integer.parseInt(stringTokens[1]));
context.write(keyValue, e);
}
}
public static class MatriceMapperN extends Mapper<Object,Text,IntWritable,Element> {
@Override
public void map(Object key, Text value, Context context)
throws IOException, InterruptedException {
String readLine = value.toString();
String[] stringTokens = readLine.split(",");
int index = Integer.parseInt(stringTokens[1]);
double elementValue = Double.parseDouble(stringTokens[2]);
Element e = new Element(1,index, elementValue);
IntWritable keyValue = new IntWritable(Integer.parseInt(stringTokens[0]));
context.write(keyValue, e);
}
}

```

```

public static class ReducerMxN extends Reducer<IntWritable,Element, Pair, DoubleWritable> {
    @Override
    public void reduce(IntWritable key, Iterable<Element> values, Context context) throws
    IOException, InterruptedException {
        ArrayList<Element> M = new ArrayList<Element>();
        ArrayList<Element> N = new ArrayList<Element>();
        Configuration conf = context.getConfiguration();
        for(Element element : values) {
            Element tempElement = ReflectionUtils.newInstance(Element.class, conf);
            ReflectionUtils.copy(conf, element, tempElement);
            if (tempElement.tag == 0) {
                M.add(tempElement);
            } else if(tempElement.tag == 1) {
                N.add(tempElement);
            }
        }
        for(int i=0;i<M.size();i++) {
            for(int j=0;j<N.size();j++) {
                Pair p = new Pair(M.get(i).index,N.get(j).index);
                double multiplyOutput = M.get(i).value * N.get(j).value;
                context.write(p, new DoubleWritable(multiplyOutput));
            }
        }
    }
}

public static class MapMxN extends Mapper<Object, Text, Pair, DoubleWritable> {
    @Override
    public void map(Object key, Text value, Context context)
    throws IOException, InterruptedException {
        String readLine = value.toString();
        String[] pairValue = readLine.split(" ");
        Pair p = new Pair(Integer.parseInt(pairValue[0]),Integer.parseInt(pairValue[1]));
        DoubleWritable val = new DoubleWritable(Double.parseDouble(pairValue[2]));
        context.write(p, val);
    }
}

public static class ReduceMxN extends Reducer<Pair, DoubleWritable, Pair,
DoubleWritable> {
    @Override
    public void reduce(Pair key, Iterable<DoubleWritable> values, Context context) throws
    IOException, InterruptedException {
        double sum = 0.0;
        for(DoubleWritable value : values) { sum += value.get();
        }
        context.write(key, new DoubleWritable(sum));
    }
}

```

```

}
public static void main(String[] args) throws Exception {
    Job job = Job.getInstance();
    job.setJobName("MapIntermediate");
    job.setJarByClass(Project1.class);
    MultipleInputs.addInputPath(job, new Path(args[0]), TextInputFormat.class,
    MatriceMapperM.class);
    MultipleInputs.addInputPath(job, new Path(args[1]), TextInputFormat.class,
    MatriceMapperN.class);
    job.setReducerClass(ReducerMxN.class);
    job.setMapOutputKeyClass(IntWritable.class);
    job.setMapOutputValueClass(Element.class);
    job.setOutputKeyClass(Pair.class);
    job.setOutputValueClass(DoubleWritable.class);
    job.setOutputFormatClass(TextOutputFormat.class);
    FileOutputFormat.setOutputPath(job, new Path(args[2]));
    job.waitForCompletion(true);
    Job job2 = Job.getInstance();
    job2.setJobName("MapFinalOutput");
    job2.setJarByClass(Project1.class);
    job2.setMapperClass(MapMxN.class);
    job2.setReducerClass(ReducerMxN.class);
    job2.setMapOutputKeyClass(Pair.class);
    job2.setMapOutputValueClass(DoubleWritable.class);
    job2.setOutputKeyClass(Pair.class);
    job2.setOutputValueClass(DoubleWritable.class);
    job2.setInputFormatClass(TextInputFormat.class);
    job2.setOutputFormatClass(TextOutputFormat.class);
    FileInputFormat.setInputPaths(job2, new Path(args[2]));
    FileOutputFormat.setOutputPath(job2, new Path(args[3])); job2.waitForCompletion(true);
}
}

```

Step 5. Compiling the program in particular folder named as operation

```

#!/bin/bash
rm -rf multiply.jar classes
module load hadoop/2.6.0
mkdir -p classes
javac -d classes -cp classes:`$HADOOP_HOME/bin/hadoop classpath` Multiply.java
jar cf multiply.jar -C classes .
echo "end"

```

Step 6. Running the program in particular folder named as operation

```

export HADOOP_CONF_DIR=/home/$USER/cometcluster
module load hadoop/2.6.0
myhadoop-configure.sh
start-dfs.sh

```

```
start-yarn.sh
hdfs dfs -mkdir -p /user/$USER
hdfs dfs -put M-matrix-large.txt /user/$USER/M-matrix-large.txt
hdfs dfs -put N-matrix-large.txt /user/$USER/N-matrix-large.txt
hadoop jar multiply.jar edu.uta.cse6331.Multiply /user/$USER/M-matrix-large.txt /user/$USER/N-matrix-large.txt /user/$USER/intermediate /user/$USER/output
rm -rf output-distr
mkdir output-distr
hdfs dfs -get /user/$USER/output/part* output-distr
stop-yarn.sh
stop-dfs.sh
myhadoop-cleanup.sh
```

Output:

```
module load hadoop/2.6.0
rm -rf output intermediate
hadoop --config $HOME jar multiply.jar edu.uta.cse6331.Multiply M-matrix-small.txt N-matrix-small.txt intermediate output.
```

Result:

EXP.NO:03

DATE:

MapReduce program to find the grades of student's

AIM:

To Develop a MapReduce program to find the grades of student's.

PROGRAMME:

```
import java.util.Scanner;
public class JavaExample
{
    public static void main(String args[])
    {
        int marks[] = new int[6];
        int i;
        float total=0, avg;
        Scanner scanner = new Scanner(System.in);
        for(i=0; i<6; i++) {
            System.out.print("Enter Marks of Subject" +(i+1) + ":");
            marks[i] = scanner.nextInt();
            total = total + marks[i];
        }
        scanner.close();
        //Calculating average here avg = total/6;
        System.out.print("The student Grade is: ");
        if(avg>=80)
        {
            System.out.print("A");
        }
        else if(avg>=60 && avg<80)
        {
            System.out.print("B");
        }
        else if(avg>=40 && avg<60)
        {
            System.out.print("C");
        }
        else
        {
            System.out.print("D");
        }
    }
}
```

OUTPUT:

Enter Marks of Subject1:40
Enter Marks of Subject2:80
Enter Marks of Subject3:80
Enter Marks of Subject4:40
Enter Marks of Subject5:60
Enter Marks of Subject6:60
The student Grade is: B

Result:

EXP.NO:04

DATE: **MapReduce program to calculate the frequency**

AIM:

To Develop a MapReduce program to calculate the frequency of a given word in agiven file

Map Function – It takes a set of data and converts it into another set of data, where individual elements are broken down into tuples (Key-Value pair).

Example – (Map function in Word Count)

Input

Set of data

Bus, Car, bus, car, train, car, bus, car, train, bus, TRAIN,BUS, buS, caR, CAR, car, BUS, TRAIN

Output

Convert into another set of data

(Key, Value)

(Bus,1), (Car,1), (bus,1), (car,1), (train,1), (car,1), (bus,1), (car,1), (train,1), (bus,1),

(TRAIN,1),(BUS,1), (buS,1), (caR,1), (CAR,1), (car,1), (BUS,1), (TRAIN,1)

Reduce Function – Takes the output from Map as an input and combines those data tuples into a smaller set of tuples.

Example – (Reduce function in Word Count)

Input Set of Tuples

(output of Map function)

(Bus,1), (Car,1), (bus,1), (car,1), (train,1), (car,1), (bus,1), (car,1), (train,1), (bus,1),

(TRAIN,1),(BUS,1),

(buS,1),(caR,1),(CAR,1), (car,1), (BUS,1), (TRAIN,1)

Output Converts into smaller set of tuples

(BUS,7), (CAR,7), (TRAIN,4)

Work Flow of Program

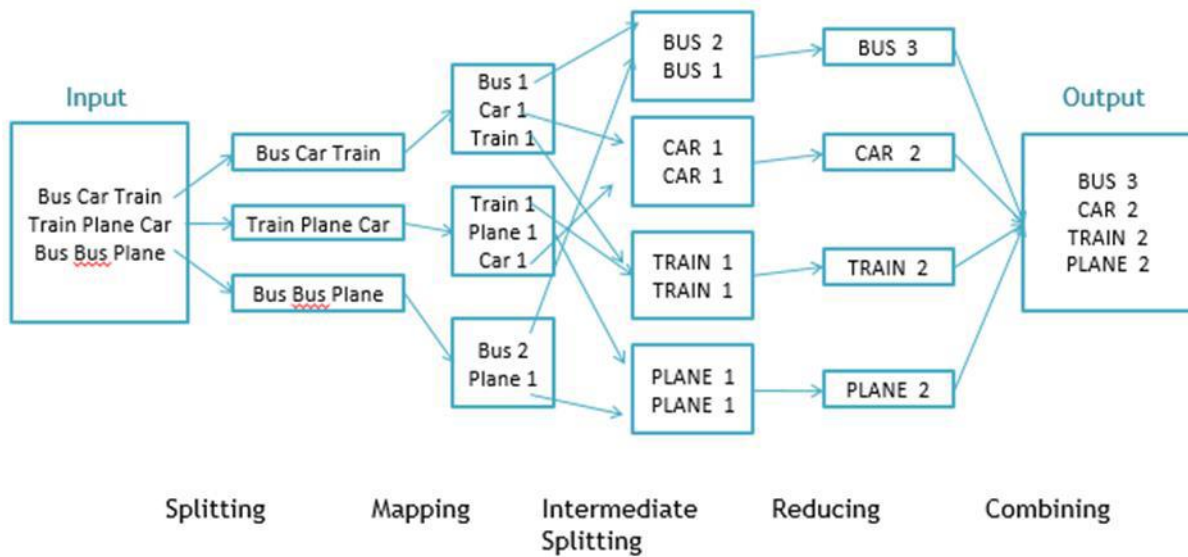


Fig. Workflow of MapReducing

Workflow of MapReduce consists of 5 steps

1. **Splitting** – The splitting parameter can be anything, e.g. splitting by space, comma, semicolon, or even by a new line (“\n”).
2. **Mapping** – as explained above
3. **Intermediate splitting** – the entire process in parallel on different clusters. In order to group them in “Reduce Phase” the similar KEY data should be on same cluster.
4. **Reduce** – it is nothing but mostly group by phase
5. **Combining** – The last phase where all the data (individual result set from each cluster) is combine together to form a Result

Now Let's See the Word Count Program in Java

Make sure that Hadoop is installed on your system with java idk

Steps to follow

Step 1. Open Eclipse> File > New > Java Project > (Name it – MRProgramsDemo) > Finish

Step 2. Right Click > New > Package (Name it - PackageDemo) > Finish

Step 3. Right Click on Package > New > Class (Name it - WordCount)

Stp 4. Add Following Reference Libraries

Right Click on Project > Build Path> Add External Archivals

- ☐ /usr/lib/hadoop-0.20/hadoop-core.jar
- ☐ Usr/lib/hadoop-0.20/lib/Commons-cli-1.2.jar

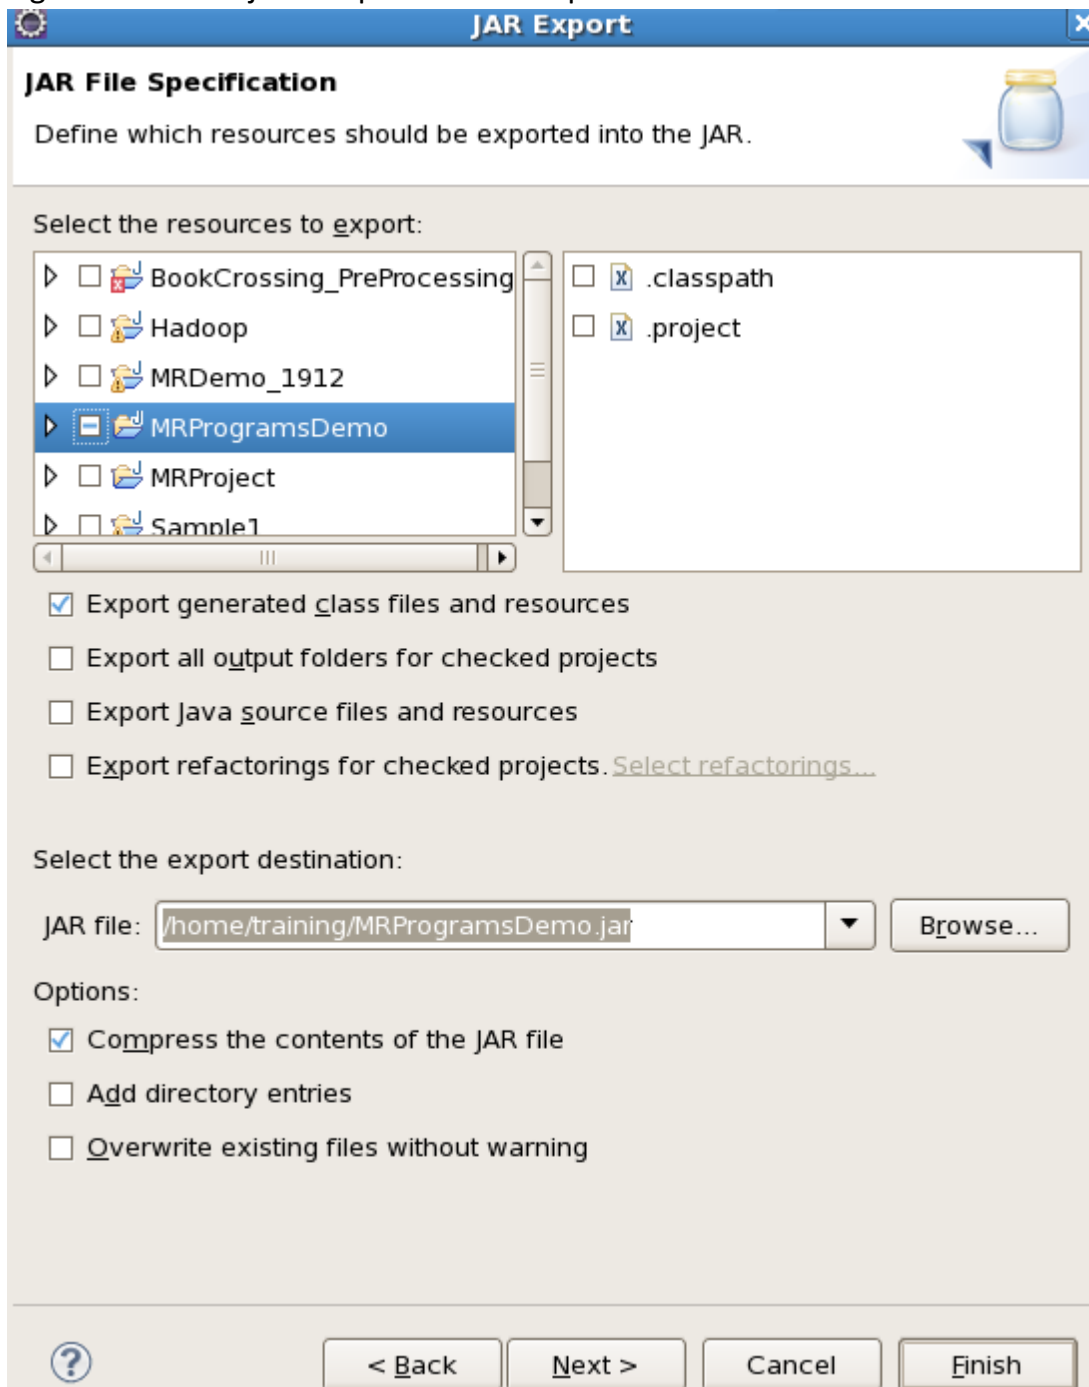
Program: Step 5. Type following Program :

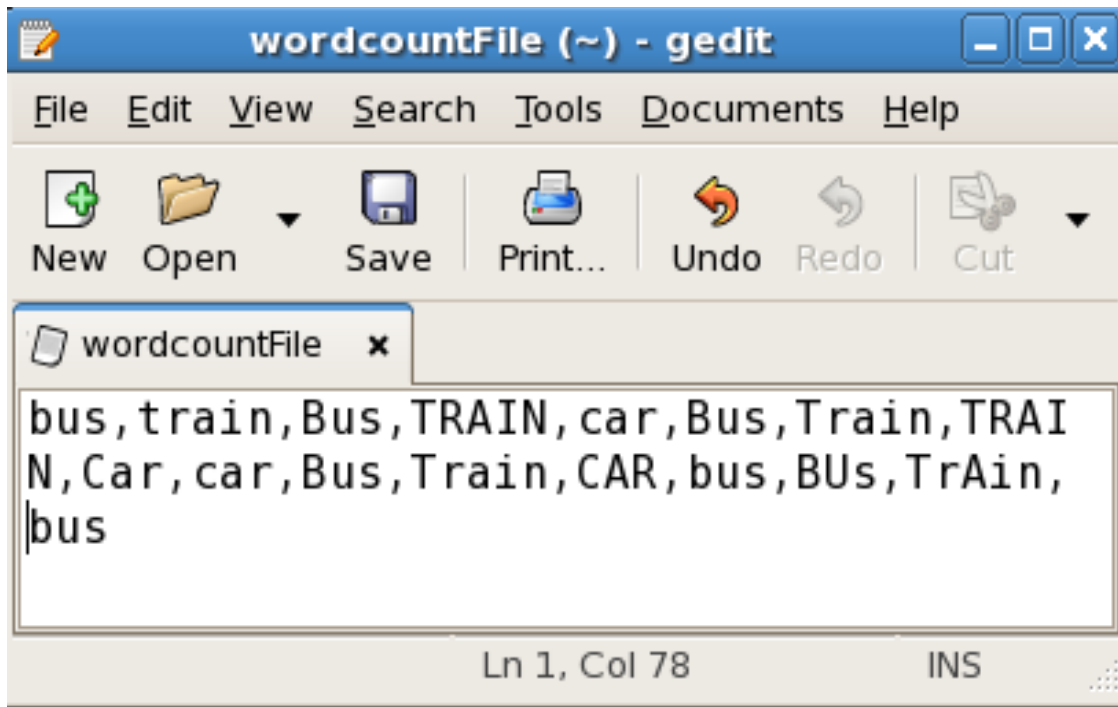
```
package PackageDemo;
import java.io.IOException;
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
import org.apache.hadoop.util.GenericOptionsParser;
public class WordCount {
public static void main(String [] args) throws Exception
{
Configuration c=new Configuration();
String[] files=new GenericOptionsParser(c,args).getRemainingArgs();
Path input=new Path(files[0]);
Path output=new Path(files[1]);
Job j=new Job(c,"wordcount");
j.setJarByClass(WordCount.class);
j.setMapperClass(MapForWordCount.class);
j.setReducerClass(ReduceForWordCount.class);
j.setOutputKeyClass(Text.class);
j.setOutputValueClass(IntWritable.class);
FileInputFormat.addInputPath(j, input);
FileOutputFormat.setOutputPath(j, output);
System.exit(j.waitForCompletion(true)?0:1);
}
public static class MapForWordCount extends Mapper<LongWritable, Text, Text,
IntWritable>{
public void map(LongWritable key, Text value, Context con) throws IOException,
InterruptedException
{
String line = value.toString();
String[] words=line.split(" ");
for(String word: words )
{
Text outputKey = new Text(word.toUpperCase().trim());
```

```
IntWritable outputValue = new IntWritable(1);
con.write(outputKey, outputValue);
}
}
}
public static class ReduceForWordCount extends Reducer<Text, IntWritable, Text,
IntWritable>
{
public void reduce(Text word, Iterable<IntWritable> values, Context con) throws
IOException,
InterruptedException
{
int sum = 0;
for(IntWritable value : values)
{
sum += value.get();
}
con.write(word, new IntWritable(sum));
}
}
}
```

Make Jar File

Right Click on Project> Export> Select export destination as Jar File > next> Finish





To Move this into Hadoop directly, open the terminal and enter the following commands:

```
[training@localhost ~]$ hadoop fs -put wordcountFile wordCountFile
```

Run Jar file

(Hadoop jar jarfilename.jar packageName.ClassName PathToInputTextFile
PathToOutputDirectry)

```
[training@localhost ~]$ Hadoop jar MRProgramsDemo.jar  
PackageDemo.WordCount wordCountFile MRDir1
```

Output:

```
[training@localhost ~]$ hadoop fs -ls MRDir1
```

Found 3 items

```
-rw-r--r-- 1 training supergroup
```

```
0 2016-02-23 03:36 /user/training/MRDir1/_SUCCESS
```

```
drwxr-xr-x - training supergroup
```

```
0 2016-02-23 03:36 /user/training/MRDir1/_logs
```

```
-rw-r--r-- 1 training supergroup
```

```
20 2016-02-23 03:36 /user/training/MRDir1/part-r-00000
```

```
[training@localhost ~]$ hadoop fs -cat MRDir1/part-r-00000
```

```
BUS 7
```

```
CAR 4
```

```
TRAIN 6
```

Result:

EXP.NO:05

DATE: MapReduce to find the maximum electrical consumption in each year

AIM:

To Develop a MapReduce to find the maximum electrical consumption in each year given electrical consumption for each month in each year.

Given below is the data regarding the electrical consumption of an organization. It contains the monthly electrical consumption and the annual average for various years. If the above data is given as input, we have to write applications to process it and produce results such as finding the year of maximum usage, year of minimum usage, and so on. This is a walkover for the programmers with finite number of records. They will simply write the logic to produce the required output, and pass the data to the application written.

But, think of the data representing the electrical consumption of all the largescale industries of a particular state, since its formation.

When we write applications to process such bulk data,

- They will take a lot of time to execute.
- There will be a heavy network traffic when we move data from source to network server and so on.

To solve these problems, we have the MapReduce framework

Input Data

The above data is saved as sample.txt and given as input. The input file looks as shown below.

```
1979 23 23 2 43 24 25 26 26 26 26 25 26 25
1980 26 27 28 28 28 30 31 31 31 30 30 30 29
1981 31 32 32 32 33 34 35 36 36 34 34 34 34
1984 39 38 39 39 39 41 42 43 40 39 38 38 40
1985 38 39 39 39 39 41 41 41 00 40 39 39 45
```

Source code:

```
import java.util.*;
import java.io.IOException;
import java.io.IOException;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.conf.*;
import org.apache.hadoop.io.*;
import org.apache.hadoop.mapred.*;
import org.apache.hadoop.util.*;
public class ProcessUnits
{
//Mapper class
public static class E_EMapper extends MapReduceBase implements
Mapper<LongWritable, /*Input key Type */ Text, /*Input value Type*/
```

```

Text, /*Output key Type*/ IntWritable> /*Output value Type*/
{
//Map function
public void map(LongWritable key, Text value, OutputCollector<Text, IntWritable> output,
Reporter reporter) throws IOException
{
String line = value.toString(); String lasttoken = null;
StringTokenizer s = new StringTokenizer(line, "\t");
String year = s.nextToken();
while(s.hasMoreTokens())
{
lasttoken=s.nextToken();
}
int avgprice = Integer.parseInt(lasttoken);
output.collect(new Text(year), new IntWritable(avgprice));
}
}
//Reducer class
public static class E_EReduce extends MapReduceBase implements
Reducer<Text, IntWritable, Text, IntWritable >
{
//Reduce function
public void reduce( Text key, Iterator <IntWritable> values, OutputCollector<Text,
IntWritable> output, Reporter reporter) throws
IOException
{
int maxavg=30;
int val=Integer.MIN_VALUE;
while (values.hasNext())
{
if((val=values.next().get())>maxavg)
{
output.collect(key, new IntWritable(val));
}
}
}
}
//Main function
public static void main(String args[])throws Exception
{
JobConf conf = new JobConf(ProcessUnits.class);
conf.setJobName("max_electricityunits");
conf.setOutputKeyClass(Text.class);
conf.setOutputValueClass(IntWritable.class);
conf.setMapperClass(E_EMapper.class);
conf.setCombinerClass(E_EReduce.class);
conf.setReducerClass(E_EReduce.class);

```

```
conf.setInputFormat(TextInputFormat.class);  
conf.setOutputFormat(TextOutputFormat.class);  
FileInputFormat.setInputPaths(conf, new Path(args[0])); FileOutputFormat.setOutputPath(conf, new Path(args[1]));  
JobClient.runJob(conf);  
}
```

Output:

Input:

Kolkata,56

Jaipur,45

Delhi,43

Mumbai,34

Goa,45

Kolkata,35

Jaipur,34

Delhi,32

Output:

Kolkata 56

Jaipur 45

Delhi 43

Mumbai 34

Result:

EXP.NO:06

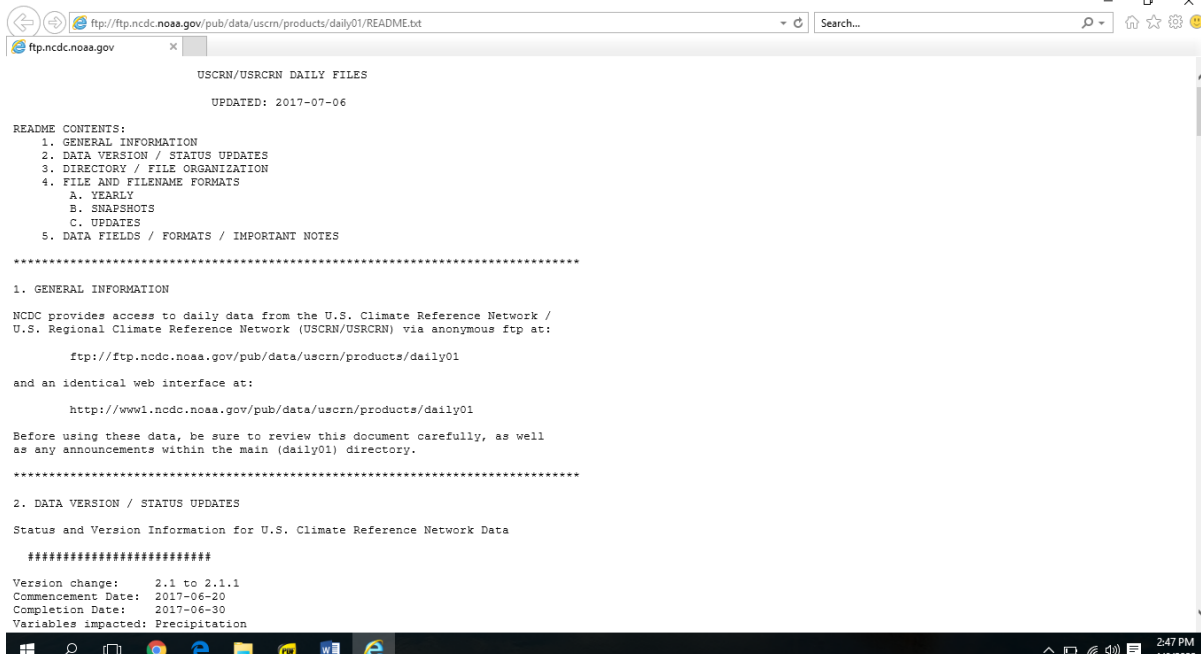
DATE: **MapReduce to analyze weather data set and print whether the day is shinny or cool**

AIM:

To Develop a MapReduce to analyze weather data set and print whether the day is shinny or cool day.

NOAA's National Climatic Data Center (**NCDC**) is responsible for preserving, monitoring, assessing, and providing public access to weather data.

NCDC provides access to daily data from the U.S. Climate Reference Network / U.S. Regional Climate Reference Network (USCRN/USRCRN) via anonymous ftp at:



```

ftp://ftp.ncdc.noaa.gov/pub/data/uscrn/products/daily01/README.txt
USCRN/USRCRN DAILY FILES
UPDATED: 2017-07-06

README CONTENTS:
1. GENERAL INFORMATION
2. DATA VERSION / STATUS UPDATES
3. DIRECTORY / FILE ORGANIZATION
4. FILE AND FILENAME FORMATS
  A. YEARLY
  B. SNAPSHOTS
  C. UPDATES
5. DATA FIELDS / FORMATS / IMPORTANT NOTES

*****

1. GENERAL INFORMATION

NCDC provides access to daily data from the U.S. Climate Reference Network /
U.S. Regional Climate Reference Network (USCRN/USRCRN) via anonymous ftp at:

    ftp://ftp.ncdc.noaa.gov/pub/data/uscrn/products/daily01

and an identical web interface at:

    http://www1.ncdc.noaa.gov/pub/data/uscrn/products/daily01

Before using these data, be sure to review this document carefully, as well
as any announcements within the main (daily01) directory.

*****

2. DATA VERSION / STATUS UPDATES

Status and Version Information for U.S. Climate Reference Network Data

*****

Version change:    2.1 to 2.1.1
Commencement Date: 2017-06-20
Completion Date:   2017-06-30
Variables impacted: Precipitation
  
```

Dataset ftp:/

After going through wordcount mapreduce guide, you now have the basic idea of how a mapreduce program works. So, let us see a complex mapreduce program on weather dataset. Here I am using one of the dataset of year 2015 of Austin, Texas . We will do analytics on the dataset and classify whether it was a hot day or a cold day NCDC gives us all the weather data we need for this mapreduce project. The dataset which we will be using looks like below snapshot.

ftp://ftp.ncdc.noaa.gov/pub/data/uscrn/products/daily01/2015/CRND0103-2015-TX_Austin_33_NW.txt depending on the temperature recorded by NCDC.

← → ↺ ftp://ftp.ncdc.noaa.gov/pub/data/uscm/products/daily01/2015/CRND0103-2015-TX_Austin_33_NW.txt 🔍 ☆ ⏹													
23907	20150101	2.423	-98.08	30.62	2.2	-0.6	0.8	0.9	6.2	1.47	C	3.7	1.1
2.5	99.9	85.4	97.2	0.369	0.308	-99.000	-99.000	-99.000	7.0	8.1	-9999.0	-9999.0	-9999.0
23907	20150102	2.423	-98.08	30.62	3.5	1.3	2.4	2.2	9.0	1.43	C	4.9	2.3
3.1	100.0	98.8	99.8	0.391	0.327	-99.000	-99.000	-99.000	7.1	7.9	-9999.0	-9999.0	-9999.0
23907	20150103	2.423	-98.08	30.62	15.9	2.3	9.1	7.5	2.9	11.00	C	16.4	2.9
7.3	100.0	34.8	73.7	0.450	0.397	-99.000	-99.000	-99.000	7.6	7.9	-9999.0	-9999.0	-9999.0
23907	20150104	2.423	-98.08	30.62	9.2	-1.3	3.9	4.2	0.0	13.24	C	12.4	-0.5
4.9	82.0	40.6	61.7	0.414	0.352	-99.000	-99.000	-99.000	7.3	7.9	-9999.0	-9999.0	-9999.0
23907	20150105	2.423	-98.08	30.62	10.9	-3.7	3.6	2.6	0.0	13.37	C	14.7	-3.0
3.8	77.9	33.3	57.4	0.399	0.340	-99.000	-99.000	-99.000	6.3	7.0	-9999.0	-9999.0	-9999.0
23907	20150106	2.423	-98.08	30.62	20.2	2.9	11.6	10.9	0.0	12.90	C	22.0	1.6
9.9	67.7	30.2	49.3	0.395	0.335	-99.000	-99.000	-99.000	8.0	8.0	-9999.0	-9999.0	-9999.0
23907	20150107	2.423	-98.08	30.62	10.9	-3.4	3.8	4.5	0.0	12.68	C	12.4	-2.1
5.5	82.7	36.5	55.7	0.387	0.328	-99.000	-99.000	-99.000	7.6	8.3	-9999.0	-9999.0	-9999.0
23907	20150108	2.423	-98.08	30.62	0.6	-7.9	-3.6	-3.3	0.0	4.98	C	3.9	-4.8
-0.5	57.7	37.6	48.1	0.372	0.316	-99.000	-99.000	-99.000	4.7	6.1	-9999.0	-9999.0	-9999.0
23907	20150109	2.423	-98.08	30.62	2.0	0.1	1.0	0.8	0.0	2.52	C	4.1	1.2
2.5	87.8	48.9	64.4	0.368	0.312	-99.000	-99.000	-99.000	5.4	6.2	-9999.0	-9999.0	-9999.0
23907	20150110	2.423	-98.08	30.62	0.5	-2.0	-0.8	-0.6	3.3	2.11	C	2.5	-0.1

PROGRAM:

```
import java.io.IOException;
import java.util.Iterator;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
import org.apache.hadoop.mapreduce.lib.output.TextOutputFormat;
```

```
import org.apache.hadoop.mapreduce.lib.input.TextInputFormat;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.conf.Configuration;
public class MyMaxMin {
public static class MaxTemperatureMapper extends
Mapper<LongWritable, Text, Text, Text> {
/**
* @method map
* This method takes the input as text data type
* Now leaving the first five tokens, it takes 6th token is taken as temp_max and
* 7th token is taken as temp_min. Now temp_max > 35 and
temp_min < 10 are passed to the reducer.
*/ @Override
public void map(LongWritable arg0, Text Value, Context context) throws IOException,
InterruptedException {
//Converting the record (single line) to String and storing it in a String variable line
String line = Value.toString();
//Checking if the line is not empty
if (!(line.length() == 0)) {
//date
String date = line.substring(6, 14);
//maximum temperature
float temp_Max = Float
parseFloat(line.substring(39, 45).trim());
//minimum temperature
float temp_Min = Float
parseFloat(line.substring(47, 53).trim());
//if maximum temperature is greater than 35 , its a hot day
```

```
if (temp_Max > 35.0) {  
    // Hot day  
    context.write(new Text("Hot Day " + date),  
        new Text(String.valueOf(temp_Max)));  
}  
//if minimum temperature is less than 10, it's a cold day  
if (temp_Min < 10) {  
    // Cold day  
    context.write(new Text("Cold Day " + date),  
        new Text(String.valueOf(temp_Min)));  
}  
}  
}  
}  
//Reducer  
/*MaxTemperatureReducer class is static and extends Reducer abstract  
having four hadoop generics type Text, Text, Text, Text.  
*/  
public static class MaxTemperatureReducer extends Reducer<Text, Text, Text, Text> {  
    public void reduce (Text Key, Iterator<Text> Values, Context context) throws IOException,  
        InterruptedException {  
        String temperature = Values.next().toString();  
        context.write(Key, new Text(temperature));  
    }  
}  
public static void main(String[] args) throws Exception {  
    Configuration conf = new Configuration();
```

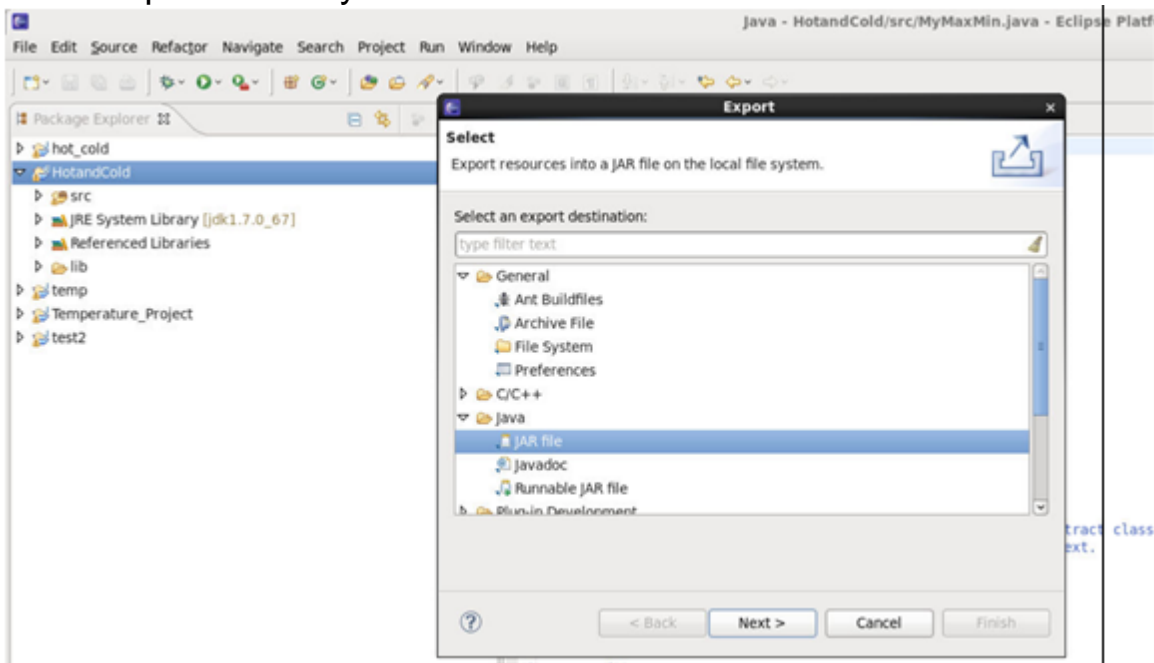
```


Job job = new Job(conf, "weather example");
job.setJarByClass(MyMaxMin.class);
job.setMapOutputKeyClass(Text.class);
job.setMapOutputValueClass(Text.class);
job.setMapperClass(MaxTemperatureMapper.class);
job.setReducerClass(MaxTemperatureReducer.class);
job.setInputFormatClass(TextInputFormat.class);
job.setOutputFormatClass(TextOutputFormat.class);
Path outputPath = new Path(args[1]);
FileInputFormat.addInputPath(job, new Path(args[0]));
FileOutputFormat.setOutputPath(job, new Path(args[1]));
OutputPath.getFileSystem(conf).delete(outputPath);
System.exit(job.waitForCompletion(true) ? 0 : 1);
}
}


```

Import the project in eclipse IDE in the same way it was told in earlier guide and change the jar paths with the jar files present in the lib directory of this project.

When the project is not having any error, we will export it as a jar file, same as we did in wordcount mapreduce guide. Right Click on the Project file and click on Export. Select jar file. Give the path where you want to save to file.



 **JAR Export** ✕

JAR Manifest Specification 

Customize the manifest file for the JAR file.

Specify the manifest:

☒ Generate the manifest file

☐ Save the manifest in the workspace

☐ Use the saved manifest in the generated JAR description file

Manifest file:

☐ Use existing manifest from workspace

Manifest file:


Seal contents:

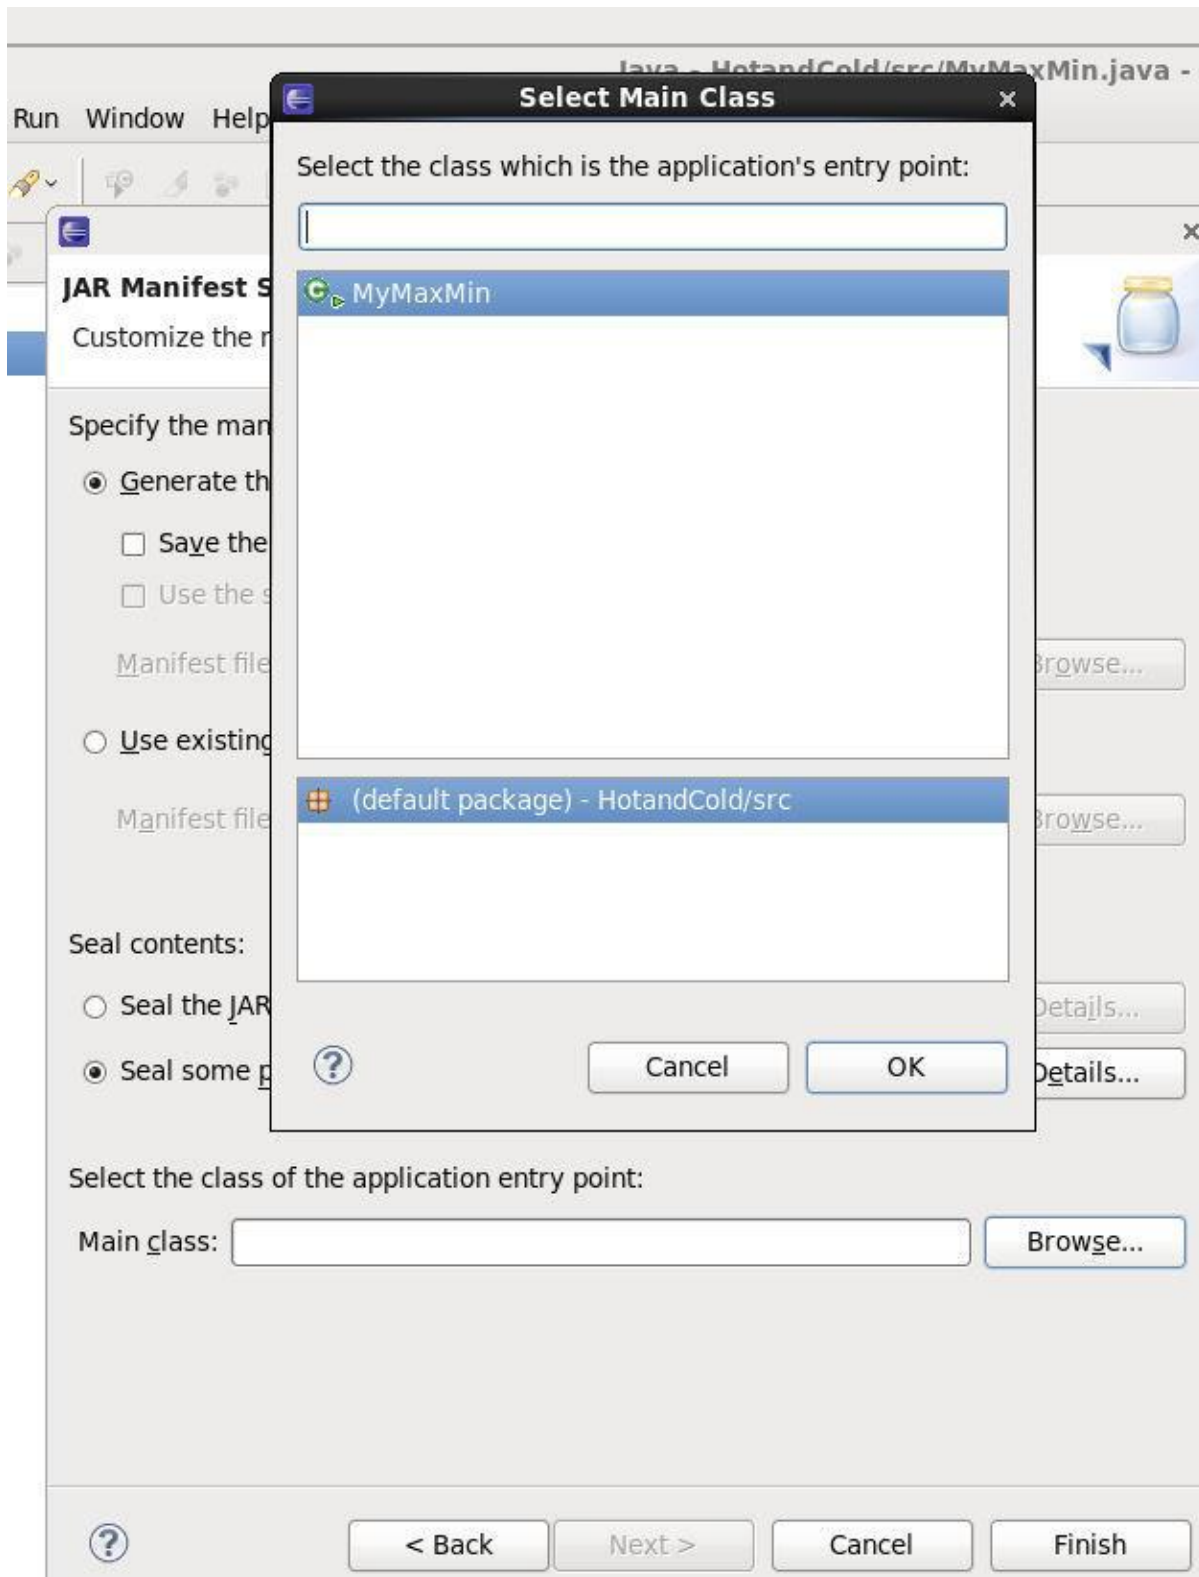
☐ Seal the JAR

☒ Seal some packages Nothing sealed

Select the class of the application entry point:

Main class:





temperature.jar

<https://drive.google.com/file/d/0B2SFMPvhXPQ5RUIZZDZSR3FYVDA/view?usp=sharing>

Download Dataset used by me using below link weather_data.txt

<https://drive.google.com/file/d/0B2SFMPvhXPQ5aFVILXAxbFh6ejA/view?usp=sharing>.

localhost:50075/browseBlock.jsp?blockId=1073742877&blockSize=1600&ger

File: [/output_hotandcold/part-r-00000](#)

Goto :

[Go back to dir listing](#)

[Advanced view/download options](#)

Cold Day	20150101	-0.6
Cold Day	20150102	1.3
Cold Day	20150103	2.3
Cold Day	20150104	-1.3
Cold Day	20150105	-3.7
Cold Day	20150106	2.9
Cold Day	20150107	-3.4
Cold Day	20150108	-7.9
Cold Day	20150109	0.1
Cold Day	20150110	-2.0
Cold Day	20150111	0.0
Cold Day	20150112	1.4
Cold Day	20150113	-0.7
Cold Day	20150114	0.9
Cold Day	20150115	1.2
Cold Day	20150116	3.5
Cold Day	20150117	5.0
Cold Day	20150118	7.6
Cold Day	20150119	6.7
Cold Day	20150120	9.5
Cold Day	20150121	6.9
Cold Day	20150122	3.5
Cold Day	20150123	2.2

Result:

DATE

XYZ.com is an online music website where users listen to various tracks

AIM:

XYZ.com is an online music website where users listen to various tracks, the data gets collected which is given below.

Write a MapReduce program to get the following

- ☐ Number of unique listeners
- ☐ Number of times the track was shared with others
- ☐ Number of times the track was listened to on the radio
- ☐ Number of times the track was listened to in total
- ☐ Number of times the track was skipped on the radio

Solution

XYZ.com is an online music website where users listen to various tracks, the data gets collected like shown below. Write a map reduce program to get following stats

- Number of unique listeners
- Number of times the track was shared with others
- Number of times the track was listened to on the radio
- Number of times the track was listened to in total
- Number of times the track was skipped on the radio

The data is coming in log files and looks like as shown below.

UserId|TrackId|Shared|Radio|Skip

111115|222|0|1|0

111113|225|1|0|0

```
111117|223|0|1|1
111115|225|1|0|0
```

In this tutorial we are going to solve the first problem, that is finding out unique listeners per track.

First of all we need to understand the data, here the first column is UserId and the second one is Track Id. So we need to write a mapper class which would emit trackId and userIds and intermediate key value pairs. To make it simple to remember the data sequence, let's create a constants class as shown below

```
public class LastFMConstants {
    public static final int USER_ID = 0; public static final int TRACK_ID = 1; public static final
    int IS_SHARED = 2; public static final int RADIO = 3;
    public static final int IS_SKIPPED = 4;
}
```

Now, let's create the mapper class which would emit intermediate key value pairs as (TrackId, UserId) as shown below

```
public static class UniqueListenersMapper extends
    Mapper< Object , Text, IntWritable, IntWritable > { IntWritable trackId = new IntWritable();
    IntWritable userId = new IntWritable();
    public void map(Object key, Text value,
        Mapper< Object , Text, IntWritable, IntWritable > .Context context)
        throws IOException, InterruptedException {
        String[] parts = value.toString().split("[|]");
        trackId.set(Integer.parseInt(parts[LastFMConstants.TRACK_ID]));
        userId.set(Integer.parseInt(parts[LastFMConstants.USER_ID]));
        if (parts.length == 5) {
            context.write(trackId, userId);
        }
    }
}
```

```
} else {  
context.getCounter(COUNTERS.INVALID_RECORD_COUNT).increment(1L)
```

```

}
}
}
public static class UniqueListenersReducer extends
Reducer< IntWritable , IntWritable, IntWritable, IntWritable> {
public void reduce( IntWritable trackId,
Iterable< IntWritable > userIds,
Reducer< IntWritable , IntWritable, IntWritable, IntWritable>.Context
context)
throws IOException, InterruptedException {
Set< Integer > userIdSet = new HashSet< Integer >();
for (IntWritable userId : userIds) {
userIdSet.add(userId.get());
}
IntWritable size = new IntWritable(userIdSet.size());
context.write(trackId, size);
}
}

```

Here we are using Set to eliminate duplicate userIds. Now we can take look at the Driver class

```

public static void main(String[] args) throws Exception { Configuration
conf = new Configuration(); if (args.length != 2) {
System.err.println("Usage: uniquelistseners < in > < out >");
System.exit(2);
}
Job job = new Job(conf, "Unique listeners per track");
job.setJarByClass(UniqueListeners.class);
job.setMapperClass(UniqueListenersMapper.class);
job.setReducerClass(UniqueListenersReducer.class);
job.setOutputKeyClass(IntWritable.class);
job.setOutputValueClass(IntWritable.class);
FileInputFormat.addInputPath(job, new Path(args[0]));
FileOutputFormat.setOutputPath(job, new Path(args[1]));
System.exit(job.waitForCompletion(true) ? 0 : 1);
org.apache.hadoop.mapreduce.Counters counters = job.getCounters();
System.out.println("No. of Invalid Records :")

```

```
+ counters.findCounter(COUNTERS.INVALID_RECORD_COUNT)
.getValue();
}
```

Output:

UserId	TrackId	Shared	Radio	Skip
--------	---------	--------	-------	------

111115	222	0	1	0
--------	-----	---	---	---

111113	225	1	0	0
--------	-----	---	---	---

111117	223	0	1	1
--------	-----	---	---	---

111115	225	1	0	0
--------	-----	---	---	---

Result:

EXP.NO:08

DATE

MapReduce program to analyze Uber data set

AIM:

To Develop a MapReduce program to analyze Uber data set to find the days on which each basement has more trips using the following dataset.

Problem Statement 1: In this problem statement, we will find the days on which each basement has more trips.

Source Code

Mapper Class:

```
public static class TokenizerMapper
extends Mapper<Object, Text, Text, IntWritable>{
java.text.SimpleDateFormat format = new
java.text.SimpleDateFormat("MM/dd/yyyy");
String[] days ={"Sun","Mon","Tue","Wed","Thu","Fri","Sat"};
private Text basement = new Text();
Date date = null;
private int trips;
public void map(Object key, Text value, Context context
) throws IOException, InterruptedException {
String line = value.toString();
String[] splits = line.split(",");
basement.set(splits[0]);
try {
date = format.parse(splits[1]);
} catch (ParseException e) {
// TODO Auto-generated catch block
e.printStackTrace();
}
trips = new Integer(splits[3]);
String keys = basement.toString()+ " "+days[date.getDay()];
context.write(new Text(keys), new IntWritable(trips));
}
}
```

Reducer Class:

```
public static class IntSumReducer
extends Reducer<Text,IntWritable,Text,IntWritable> {
private IntWritable result = new IntWritable();
public void reduce(Text key, Iterable<IntWritable> values,
Context context
) throws IOException, InterruptedException {
int sum = 0
for (IntWritable val : values) {
sum += val.get();
}
result.set(sum);
context.write(key, result);
}
}
```

Whole Source Code:

```
import java.io.IOException;
import java.text.ParseException;
import java.util.Date;
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
public class Uber1 {
public static class TokenizerMapper
extends Mapper<Object, Text, Text, IntWritable>{
java.text.SimpleDateFormat format = new
java.text.SimpleDateFormat("MM/dd/yyyy");
String[] days ={"Sun","Mon","Tue","Wed","Thu","Fri","Sat"};
private Text basement = new Text();
Date date = null;
private int trips;
public void map(Object key, Text value, Context context
```

```

) throws IOException, InterruptedException {
    String line = value.toString();
    String[] splits = line.split(",");
    basement.set(splits[0]);
    try {
        date = format.parse(splits[1]);
    } catch (ParseException e) {
        // TODO Auto-generated catch block
        e.printStackTrace();
    }
    trips = new Integer(splits[3]);
    String keys = basement.toString()+ " "+days[date.getDay()];
    context.write(new Text(keys), new IntWritable(trips));
}
}

public static class IntSumReducer
    extends Reducer<Text,IntWritable,Text,IntWritable>
    private IntWritable result = new IntWritable();
    public void reduce(Text key, Iterable<IntWritable> values
        Context context
    ) throws IOException, InterruptedException {
        int sum = 0;
        for (IntWritable val : values) {
            sum += val.get();
        }
        result.set(sum);
        context.write(key, result);
    }
}

public static void main(String[] args) throws Exception {
    Configuration conf = new Configuration();
    Job job = Job.getInstance(conf, "Uber1");
    job.setJarByClass(Uber1.class);
    job.setMapperClass(TokenizerMapper.class);
    job.setCombinerClass(IntSumReducer.class);
    job.setReducerClass(IntSumReducer.class);
    job.setOutputKeyClass(Text.class);
    job.setOutputValueClass(IntWritable.class);
}

```

```
FileInputFormat.addInputPath(job, new Path(args[0]));  
FileOutputFormat.setOutputPath(job, new Path(args[1]));  
System.exit(job.waitForCompletion(true) ? 0 : 1);  
}  
}
```

Running the Program:

First, we need to build a jar file for the above program and we need to run it as a normal Hadoop program by passing the input dataset and the output file path as shown below.

```
hadoop jar uber1.jar /uber /user/output1
```

In the output file directory, a part of the file is created and contains the below

Output:

```
B02512 Sat 15026  
B02512 Sun 10487  
B02512 Thu 15809  
B02512 Tue 12041  
B02512 Wed 12691  
B02598 Fri 93126  
B02598 Mon 60882  
B02598 Sat 94588  
B02598 Sun 66477  
B02598 Thu 90333  
B02598 Tue 63429  
B02598 Wed 71956  
B02617 Fri 125067  
B02617 Mon 80591  
B02617 Sat 127902  
B02617 Sun 91722  
B02617 Thu 118254  
B02617 Tue 86602  
B02617 Wed 94887  
B02682 Fri 114662  
B02682 Mon 74939  
B02682 Sat 120283  
B02682 Sun 82825  
B02682 Thu 106643  
B02682 Tue 76905  
B02682 Wed 86252
```

B02764 Fri 326968
B02764 Mon 214116
B02764 Sat 356789
B02764 Sun 249896
B02764 Thu 304200
B02764 Tue 221343
B02764 Wed 241137
B02765 Fri 34934
B02765 Mon 21974
B02765 Sat 36737

Result:

EXP.NO:9

DATE

HIVE OPERATIONS

AIM:

To Use Hive to create, alter, and drop databases, tables, views, functions, and indexes.

RESOURCES:

VMWare, XAMPP Server, Web Browser, 1GB RAM, Hard Disk 80 GB.

PROGRAM LOGIC:

SYNTAX for HIVE Database Operations

DATABASE Creation

CREATE DATABASE|SCHEMA [IF NOT EXISTS] <database name>

Drop Database Statement

DROP DATABASE Statement DROP (DATABASE|SCHEMA) [IF EXISTS]

database_name [RESTRICT|CASCADE];

Creating and Dropping Table in HIVE

CREATE [TEMPORARY] [EXTERNAL] TABLE [IF NOT EXISTS]

[db_name.]

table_name

[(col_name data_type [COMMENT col_comment], ...)]

[COMMENT table_comment] [ROW FORMAT row_format] [STORED AS file_format]

Loading Data into table log_data

Syntax:

LOAD DATA LOCAL INPATH '<path>/u.data' OVERWRITE INTO TABLE

u_data;

Alter Table in HIVE

Syntax

ALTER TABLE name RENAME TO new_name

ALTER TABLE name ADD COLUMNS (col_spec[, col_spec ...])

ALTER TABLE name DROP [COLUMN] column_name

ALTER TABLE name CHANGE column_name new_name new_type

ALTER TABLE name REPLACE COLUMNS (col_spec[, col_spec ...])

Creating and Dropping View

CREATE VIEW [IF NOT EXISTS] view_name [(column_name

[COMMENT

column_comment], ...)] [COMMENT table_comment] AS SELECT ...

Dropping View

Syntax:

DROP VIEW view_name

Functions in HIVE

String Functions:- round(), ceil(), substr(), upper(), reg_exp() etc Date

and Time Functions:- year(), month(), day(), to_date() etc

Aggregate Functions :- sum(), min(), max(), count(), avg() etc

INDEXES

CREATE INDEX index_name ON TABLE base_table_name (col_name, ...)

AS 'index.handler.class.name'

[WITH DEFERRED REBUILD]

[IDXPROPERTIES (property_name=property_value, ...)]

[IN TABLE index_table_name]

[PARTITIONED BY (col_name, ...)]

[

[ROW FORMAT ...] STORED AS ...

| STORED BY ...

]

[LOCATION hdfs_path]

[TBLPROPERTIES (...)]

Creating Index

CREATE INDEX index_ip ON TABLE log_data(ip_address) AS

'org.apache.hadoop.hive.ql.index.compact.CompactIndexHandler' WITH

DEFERRED

REBUILD;

Altering and Inserting Index

ALTER INDEX index_ip_address ON log_data REBUILD;

Storing Index Data in Metastore

SET

hive.index.compact.file=/home/administrator/Desktop/big/metastore_db/t

mp/index_ipadd

ress_result;

SET

hive.input.format=org.apache.hadoop.hive.ql.index.compact.HiveCompa

ctIndexInputFor

mat;

Dropping Index

DROP INDEX INDEX_NAME on TABLE_NAME;

Output:

Result:

EXP.NO:10

DATE

Queries to sort and aggregate the data in a table using HiveQL

AIM:

To Write queries to sort and aggregate the data in a table using HiveQL.

Description:

Hive is an open-source data warehousing solution built on top of Hadoop. It supports an SQL-like query language called HiveQL. These queries are compiled into MapReduce jobs that are executed on Hadoop. While Hive uses Hadoop for execution of queries, it reduces the effort that goes into writing and maintaining MapReduce jobs.

Hive supports database concepts like tables, columns, rows and partitions. Both primitive (integer, float, string) and complex data-types(map, list, struct) are supported. Moreover, these types can be composed to support structures of arbitrary complexity. The tables are serialized/deserialized using default serializers/deserializer. Any new data format and type can be supported by implementing SerDe and ObjectInspector java interface.

HiveQL - ORDER BY and SORT BY Clause

By using HiveQL ORDER BY and SORT BY clause, we can apply sort on the column. It returns the result set either in ascending or descending order. Here, we are going to execute these clauses on the records of the below table:

emp

Id	Name	Salary	Department
1	Gaurav	30000	Developer
2	Aryan	20000	Manager
3	Vishal	40000	Manager
4	John	10000	Trainer
5	Henry	25000	Developer
6	William	9000	Developer
7	Lisa	25000	Manager
8	Ronit	20000	Trainer



In HiveQL, ORDER BY clause performs a complete ordering of the query result set. Hence, the complete data is passed through a single reducer. This may take much time in the execution of large datasets. However, we can use LIMIT to minimize the sorting time.

Example:

Select the database in which we want to create a table.

hive> use hiveql;

```
codegyani@ubuntu64server: ~
hive> use hiveql;
OK
Time taken: 0.067 seconds
hive> █
```

Now, create a table by using the following command:

hive> create table emp (Id int, Name string , Salary float, Department string)

row format delimited

fields terminated by ',' ;

```
codegyani@ubuntu64server: ~
hive> create table emp (Id int, Name string , Salary float, Department string)
> row format delimited
> fields terminated by ',' ;
OK
Time taken: 0.419 seconds
hive>
```

Load the data into the table

hive> load data local inpath '/home/codegyani/hive/emp_data' into table emp;

```
codegyani@ubuntu64server: ~
hive> load data local inpath '/home/codegyani/hive/emp_data' into table emp;
Loading data to table hiveql.emp
Table hiveql.emp stats: [numFiles=1, totalSize=200]
OK
Time taken: 1.411 seconds
hive>
```

Now, fetch the data in the descending order by using the following command
hive> select * from emp order by salary desc;



```
codegyani@ubuntu64server: ~
hive> select * from emp order by salary desc;
Query ID = codegyani_20190802063522_65b28a82-4d0b-492a-ae25-2faef1471b65
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1555046592674_0032, Tracking URL = http://ubuntu64server:8088
/proxy/application_1555046592674_0032/
Kill Command = /home/codegyani/hadoop-2.7.1/bin/hadoop job -kill job_155504659
2674_0032
```

```
codegyani@ubuntu64server: ~
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2019-08-02 06:36:48,908 Stage-1 map = 0%, reduce = 0%
2019-08-02 06:37:49,829 Stage-1 map = 0%, reduce = 0%
2019-08-02 06:38:02,548 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 7.03 sec
2019-08-02 06:39:03,090 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 10.31 sec
2019-08-02 06:39:18,347 Stage-1 map = 100%, reduce = 67%, Cumulative CPU 12.98 sec
2019-08-02 06:39:35,537 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 22.34 sec
MapReduce Total cumulative CPU time: 22 seconds 340 msec
Ended Job = job_1555046592674_0032
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 22.34 sec HDFS Read: 6788 H
DFS Write: 227 SUCCESS
```

```
codegyani@ubuntu64server: ~
Total MapReduce CPU Time Spent: 22 seconds 340 msec
OK
3      "Vishal"          40000.0 Manager
1      "Gaurav"         30000.0 Developer
7      "Lisa"           25000.0 Manager
5      "Henry"          25000.0 Developer
8      "Ronit"          20000.0 Trainer
2      "Aryan"          20000.0 Manager
4      "John"           10000.0 Trainer
6      "William"        9000.0 Developer
NULL   NULL            NULL      NULL
Time taken: 257.304 seconds, Fetched: 9 row(s)
hive>
```

HiveQL - SORT BY Clause

The HiveQL SORT BY clause is an alternative of ORDER BY clause. It orders the data within each reducer. Hence, it performs the local ordering, where each reducer's output is sorted separately. It may also give a partially ordered result

Example:

Let's fetch the data in the descending order by using the following command

```
hive> select * from emp sort by salary desc;
```



```
codegyani@ubuntu64server: ~
hive> select * from emp sort by salary desc;
Query ID = codegyani_20190802065014_f877314f-8d92-428f-8a9f-1b6a9b67c328
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1555046592674_0033, Tracking URL = http://ubuntu64server:8088/
proxy/application_1555046592674_0033/
Kill Command = /home/codegyani/hadoop-2.7.1/bin/hadoop job -kill job_1555046592674_0033
```

```
codegyani@ubuntu64server: ~
Total MapReduce CPU Time Spent: 22 seconds 690 msec
OK
3      "Vishal"      40000.0 Manager
1      "Gaurav"     30000.0 Developer
7      "Lisa"       25000.0 Manager
5      "Henry"      25000.0 Developer
8      "Ronit"      20000.0 Trainer
2      "Aryan"      20000.0 Manager
4      "John"       10000.0 Trainer
6      "William"    9000.0 Developer
NULL   NULL        NULL      NULL
Time taken: 268.62 seconds, Fetched: 9 row(s)
```

Cluster By:

Cluster By used as an alternative for both Distribute BY and Sort BY clauses in Hive-QL.

Cluster BY clause used on tables present in Hive. Hive uses the columns in Cluster by to distribute the rows among reducers. Cluster BY columns will go to the multiple reducers.

- It ensures sorting orders of values present in multiple reducers

For example, Cluster By clause mentioned on the Id column name of the table employees_guru table. The output when executing this query will give results to multiple reducers at the back end. But as front end it is an alternative clause for both Sort By and Distribute By.

Example:

SELECT Id, Name from employees_guru CLUSTER BY Id;

```
hive> Select Id,Name from employees guru CLUSTER BY Id;
Query ID = h1es_20151105165000_72cedc06-a797-48b1-a120-
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks not
In order to change the ave
set hive.exec.reducers.b
In order to limit the maxim
set hive.exec.reducers.m
In order to set a constant
set mapred.reduce.tasks=<number>
Starting Job = job_201511051442_0009, Tracking URL = http
Kill Command = /usr/local/hadoop-1.2.1/libexec/./bin/had
Hadoop job information for Stage-1: number of mappers: 1;
2015-11-05 16:50:08,541 Stage-1 map = 0% reduce = 0%
2015-11-05 16:50:10,546 Stage-1 map = 0% reduce = 0%,
2015-11-05 16:50:17,563 Stage-1 map = 0% reduce = 100%
MapReduce Total cumulat
Ended Job = job_201511051442_0009
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 1.6 se
Total MapReduce CPU Time Spent: 1 seconds 600 msec
OK
101 Rajesh
102 Rajiv
103 Animesh
104 Anirudh
105 Santosh
106 Ramesh
107 Sravanthi
108 Sravan
109 Suresh
110 Ravi
111 Syam
Time taken: 18.941 seconds, Fetched: 11 row(s)
```

cluster by query

cluster by query output

2

Result:

